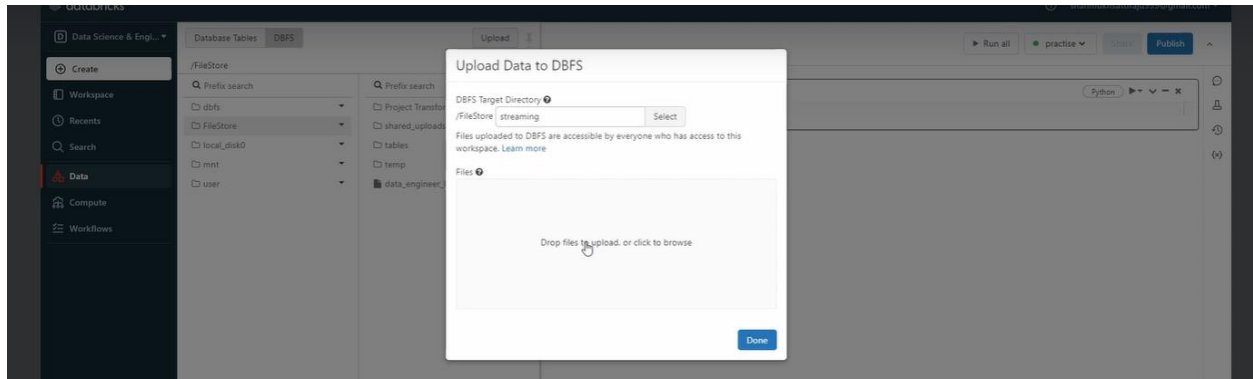
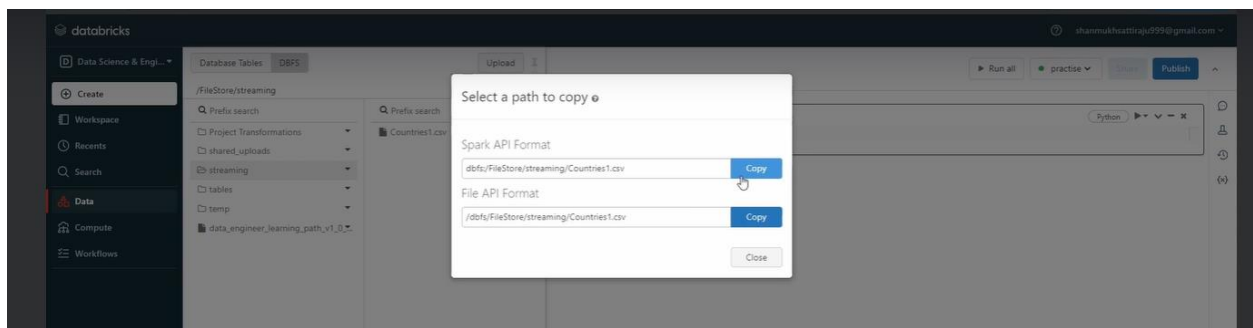
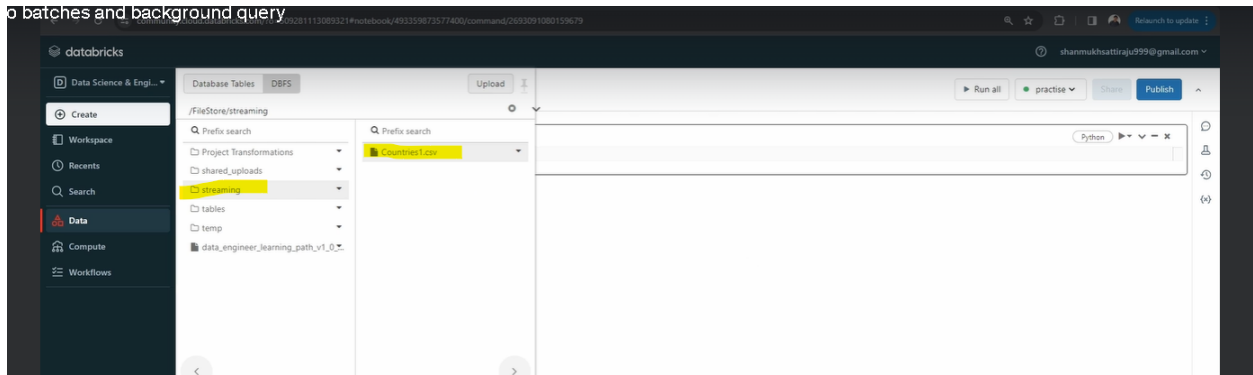


Streaming_practice

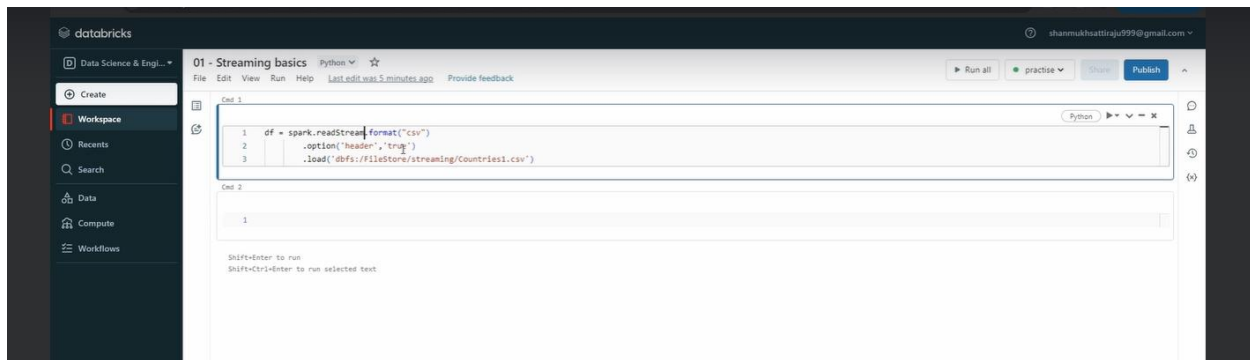


Upload a file to dbfs in the workspace creating a folder

o batches and background query

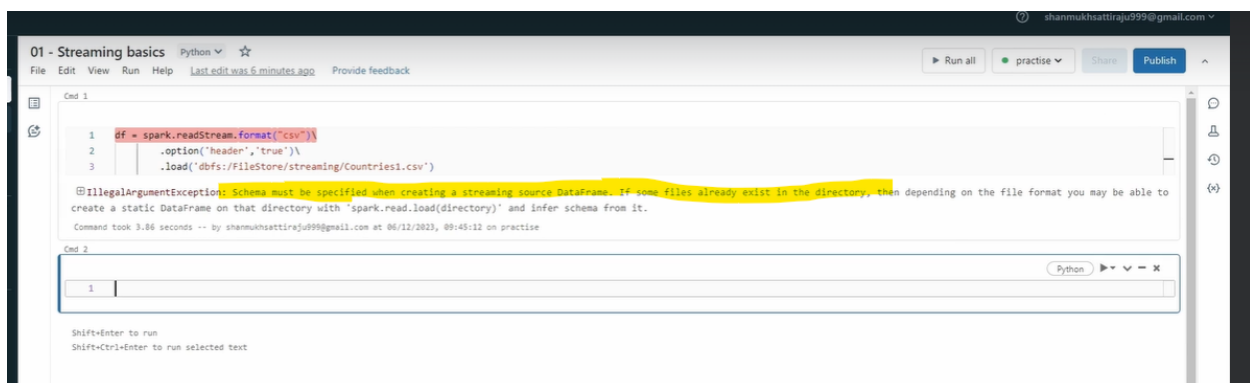


Reading the data



Reading a streaming dataset

If we try to read the dataframe directly without specifying or inferring the schema, we will get the below error.



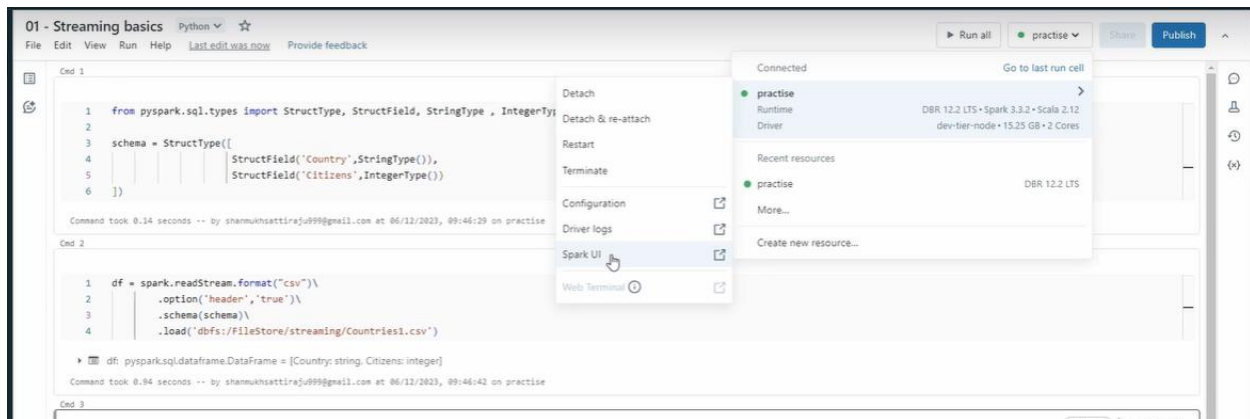
```
1 from pyspark.sql.types import StructType, StructField, StringType, IntegerType, FloatType
2
3 schema = StructType([
4     StructField('Country', StringType()),
5     StructField('Citizens', IntegerType())
6 ])

```


Define the schema and re-run the data frame



Rerunning the data frame, specifying the schema will read the stream data



Jobs initiated can be read through the spark UI interface



Initially no jobs were initiated, jobs will be initiated only when trying to fetch some data.



Streaming data can be read through display command, however basePath must be a directory.

```
Cmd 2

1 df = spark.readStream.format("csv")\
2   .option('header','true')\
3   .schema(schema)\
4   .load('dbfs:/FileStore/streaming/')

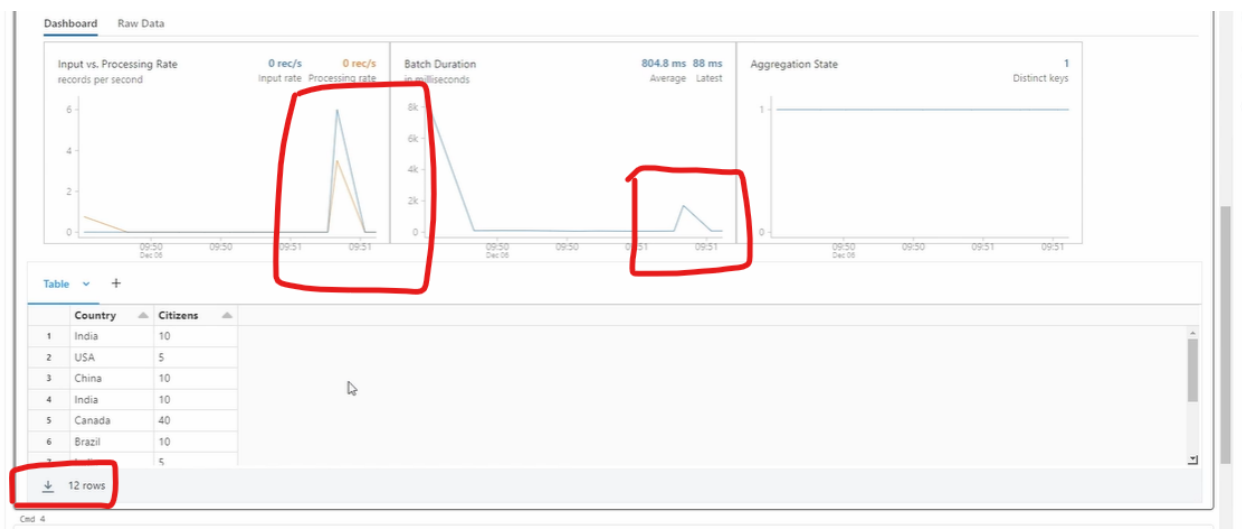
df: pyspark.sql.dataframe.DataFrame = [Country: string, Citizens: integer]
Command took 0.70 seconds -- by shanmukhsattiraju999@gmail.com at 06/12/2023, 09:49:26 on practise

Cmd 3

1 display(df)

Cancel ***
▶ (1) Spark Jobs
Stream initializing...
```

To fix it change the load path to directly to directory level.



To observe the change in the streaming input data, add a new file to the directory. Within a few seconds, it will be noticed in the input and batch streaming graph, without re-executing the display command. Record count is changed to 12 from 6.

Compute > practise

Configuration Notebooks (1) Libraries Event log **Spark UI** Driver logs Metrics Apps Spark compute UI - Master

Hostname: ec2-54-202-135-114.us-west-2.compute.amazonaws.com Spark Version: 12.2.x-scala2.12

Jobs Stages Storage Environment Executors SQL / DataFrame JDBC/ODBC Server **Structured Streaming**

Streaming Query

▼ Active Streaming Queries (1)

Page: 1 1 Pages. Jump to 1 . Show 100 items in a page. Go

Name	Status	ID	Run ID	Start Time	Duration	Avg Input /sec	Avg Process /sec	Latest Batch
display_query_2	RUNNING	26f6338a-8a1a-4d3a-8362-4f741dede6af	6d08201c-c39e-47e4-aa03-1078f9cb65fb	2023/12/06 04:19:30	6 minutes 25 seconds	0.17	0.12	2

Page: 1 1 Pages. Jump to 1 . Show 100 items in a page. Go

▼ Completed Streaming Queries (1)

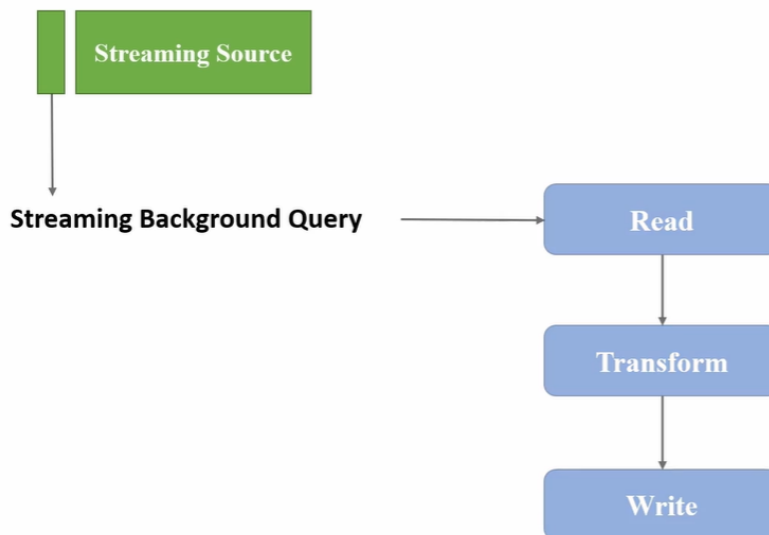
Page: 1 1 Pages. Jump to 1 . Show 100 items in a page. Go

Name	Status	ID	Run ID	Start Time	Duration	Avg Input /sec	Avg Process /sec	Latest Batch	Error
display_query_1	FAILED	2807e9a9-c9d9-4968-a165-c3f973ce5b6e	a98faeb3-e758-4a9f-b8c9-1a7e41872d5d	2023/12/06 04:18:17	0 ms	NaN	NaN	NaN	java.lang.IllegalArgumentException: Option 'basePath' must be a directory +details

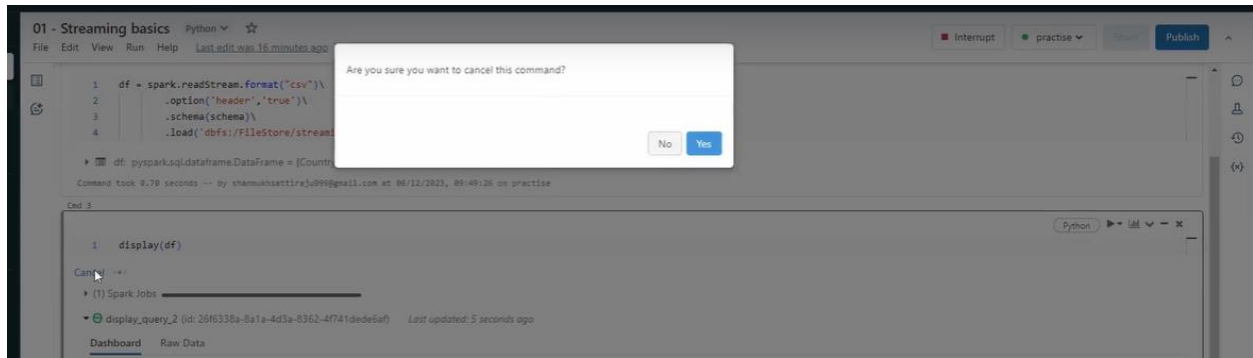
Page: 1 1 Pages. Jump to 1 . Show 100 items in a page. Go

Each file addition is read in micro batches.

Spark Structured Streaming flow



The background streaming query acts as a file watcher, whenever there is a new file added to the directory, a new job will be invoked, the file will be read, transformed and written. The query will be running continuously, until it is stopped explicitly.



Streaming queries can be stopped explicitly by cancelling the query.

Supported Sources and Sinks

Sources

File Source

Kafka Source

Socket Source

Rate Source

Table



Sinks

File Sink

Kafka Sink

Foreach Sink

Console Sink

Table

Supported sources and sinks for the Spark Streaming.

StreamWriter

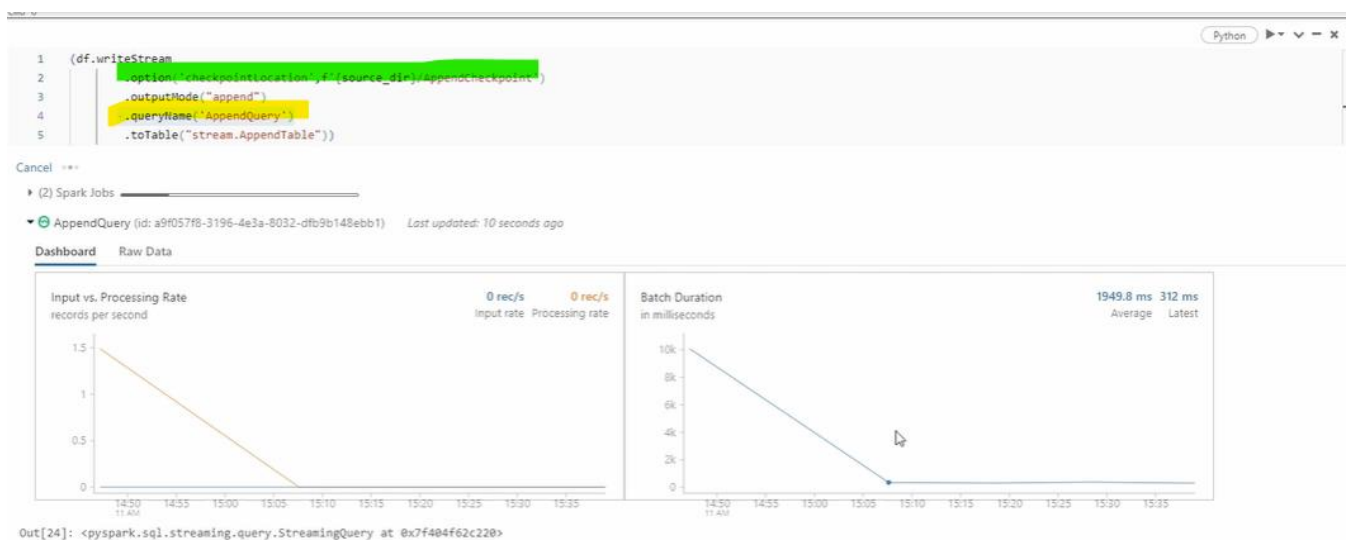
```
<StreamingDataframe>.writeStream  
    .option('checkpointLocation',<Location>)  
    .outputMode('append')  
    .toTable('<TableName>')
```

Checkpoint

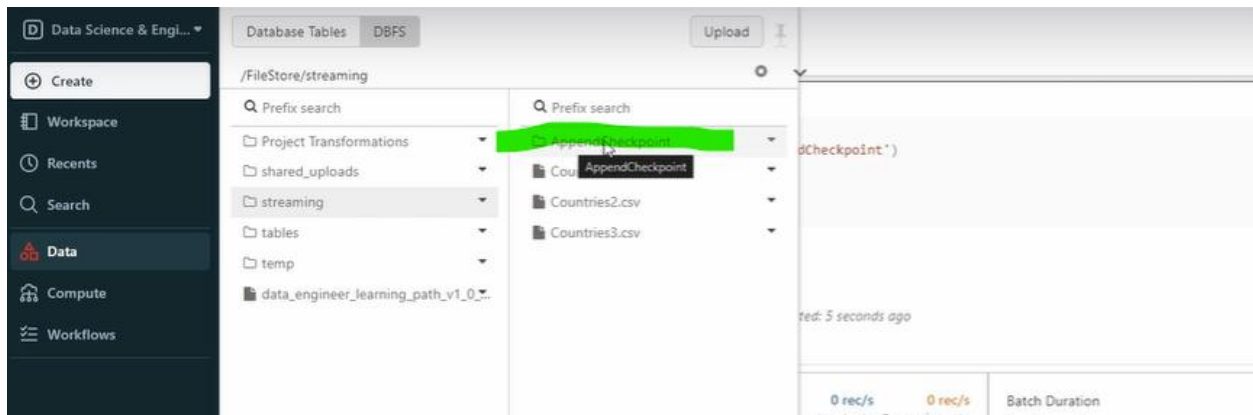
To develop fault tolerant and resilient Spark APPS

It maintains intermediate state on fault-tolerant compatible files like DBFS, ADLS and S3 storage systems to recover from the failures.

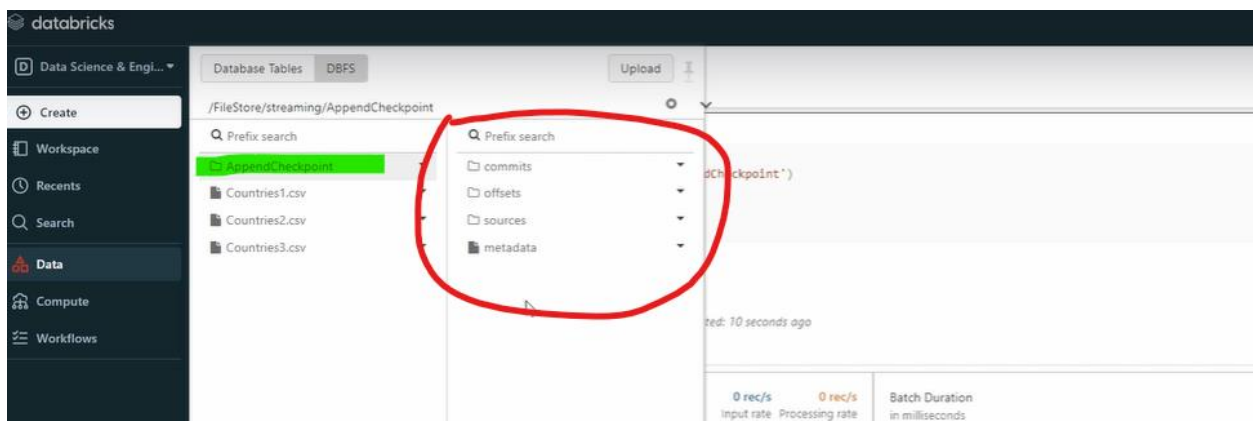
It is unique to each stream.



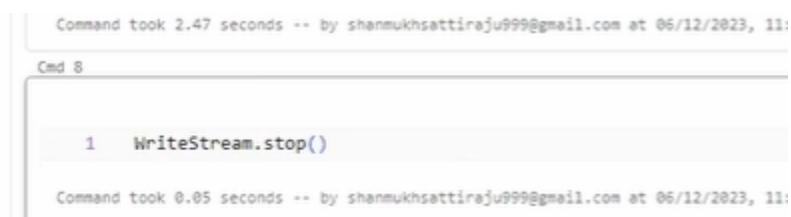
Streaming data frame, data is read, checkpoint location is specified and query name as well specified.



Now after executing the streaming query, a check point location folder is added in the location specified in the above query.



Contains the metadata of the streaming



Assign the streaming query to a variable and to stop the streaming use the command `<var>.stop()`


```
02. outputModes Python ☆
File Edit View Run Help Last edit was 19 minutes ago Provide feedback

Cmd 1
1 dbutils.fs.rm('dbfs:/user/hive/warehouse/stream.db',True)

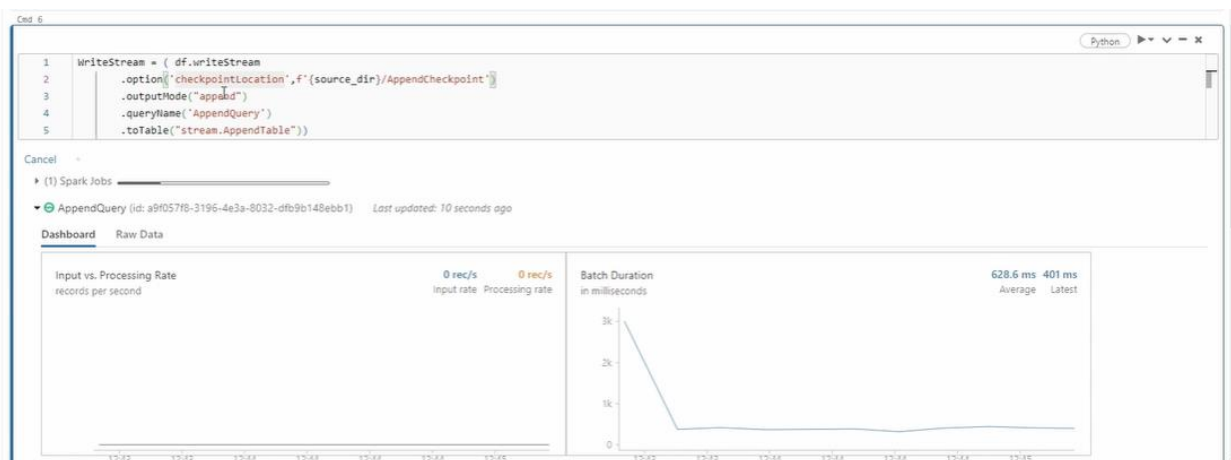
Out[12]: True
Command took 0.35 seconds -- by shanmukhsattiraju999@gmail.com at 06/12/2023, 12:26:00 on new

Cmd 2
1 %sql
2 DROP DATABASE IF EXISTS stream CASCADE;
3 CREATE DATABASE IF NOT EXISTS stream

SQL cell result stored as PySpark data frame _sqldf. Learn more
Command took 2.24 seconds -- by shanmukhsattiraju999@gmail.com at 06/12/2023, 12:26:37 on new

Cmd 3
```

In the community edition when the cluster is terminated all the schema/database created in the working session will be terminated. Therefore, while using the community edition we should remove the files in the dbfs path, drop the database if exists and re-create it.



When we are trying to run the stream query initially it will verify the checkpoint location, if there are any unprocessed files, only then it will show the updates in the graph

Output Modes

Deleting the checkpoint location will help to process the older data, without removing the data, just by removing the checkpoint, previous data will be processed.

Complete Mode

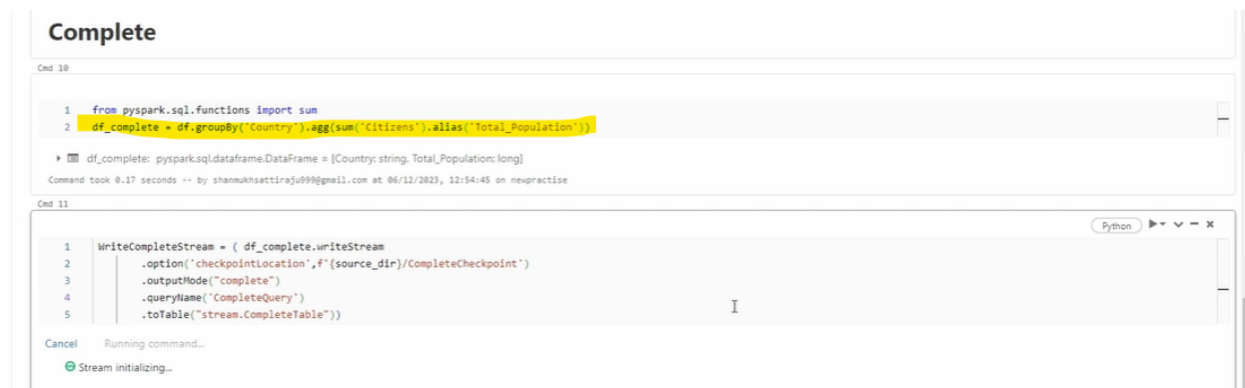


The screenshot shows a Databricks notebook interface. At the top, the tab is labeled "Complete". Below the code editor, a message indicates "(4) Spark Jobs". An "AnalysisException" is displayed with the text: "Complete output mode not supported when there are no streaming aggregations on streaming DataFrames/Datasets;". The exception details show "FileSource[dbfs:/FileStore/streaming/]" and "Command took 7.33 seconds -- by shanmukhsattiraju999@gmail.com at 06/12/2023, 12:49:33 on newpractise".

```
1 WriteCompleteStream = ( df.writeStream
2   .option('checkpointLocation',f'{source_dir}/CompleteCheckpoint')
3   .outputMode("complete")
4   .queryName("CompleteQuery")
5   .toTable("stream.CompleteTable"))
```

AnalysisException: Complete output mode not supported when there are no streaming aggregations on streaming DataFrames/Datasets;
FileSource[dbfs:/FileStore/streaming/]
Command took 7.33 seconds -- by shanmukhsattiraju999@gmail.com at 06/12/2023, 12:49:33 on newpractise

To use the complete output mode aggregation should be used on the data frame.



The screenshot shows a Databricks notebook interface. The tab is labeled "Complete". The code editor shows two lines of Python code: "from pyspark.sql.functions import sum" and "df_complete = df.groupBy('Country').agg(sum('Citizens').alias('Total_Population'))". Below the code, a message indicates "df_complete: pyspark.sql.dataframe.DataFrame = [Country: string, Total_Population: long]" and "Command took 0.17 seconds -- by shanmukhsattiraju999@gmail.com at 06/12/2023, 12:54:45 on newpractise". The bottom of the notebook shows a "WriteCompleteStream" command being executed, with a status bar indicating "Stream initializing..." and a "Cancel" button.

```
1 from pyspark.sql.functions import sum
2 df_complete = df.groupBy('Country').agg(sum('Citizens').alias('Total_Population'))
```

df_complete: pyspark.sql.dataframe.DataFrame = [Country: string, Total_Population: long]
Command took 0.17 seconds -- by shanmukhsattiraju999@gmail.com at 06/12/2023, 12:54:45 on newpractise

```
1 WriteCompleteStream = ( df_complete.writeStream
2   .option('checkpointLocation',f'{source_dir}/CompleteCheckpoint')
3   .outputMode("complete")
4   .queryName("CompleteQuery")
5   .toTable("stream.CompleteTable"))
```

Cancel Running command...
Stream initializing...

Applied aggregation to execute the **complete output mode**

Use case of complete mode

When a new file or dataset is added, and we need to apply the aggregation on top it, to fetch the results of the query entire files should be read.

	A	B	C	D	E	F	G	H	I	J	K	L
1	Countries1											
2	Country	Citizens							Result of Countries 1 and 2			
3	India	10							Country	Total Pop		
4	USA	5							India	30		
5	China	10							USA	15		
6	India	10							China	15		
7	Canada	40							Canada	50		
8	Brazil	10							Brazil	60		
9												
10	Countries2											
11	Country	Citizens										
12	India	5										
13	USA	10										
14	China	5										
15	India	5										
16	Canada	10										
17	Brazil	50										

Trigger types

01. Trigger - default or unspecified Trigger

Cmd 7

```
1 WriteStream = ( df.writeStream
2   .option('checkpointLocation',f'{source_dir}/AppendCheckpoint')
3   .outputMode("append")
4   .queryName('DefaultTrigger')
5   .toTable("stream.AppendTable"))
```

Cancel

▶ (2) Spark Jobs

If no trigger time is specified, by default for every 5sec, streaming query will look for new data. If the new file is added every 5 to 10 min, default trigger is not a good option.

Cmd 9

02. Trigger - processingTime

Cmd 10

```

1 WriteStream = df.writeStream
2   .option('checkpointLocation',f'{source_dir}/AppendCheckpoint')
3   .outputMode("append")
4   .trigger(processingTime='2 minutes')
5   .queryName('ProcessingTime')
6   .toTable("stream.AppendTable"))

```

File 11

Every 2 min, a micro batch will be processed.

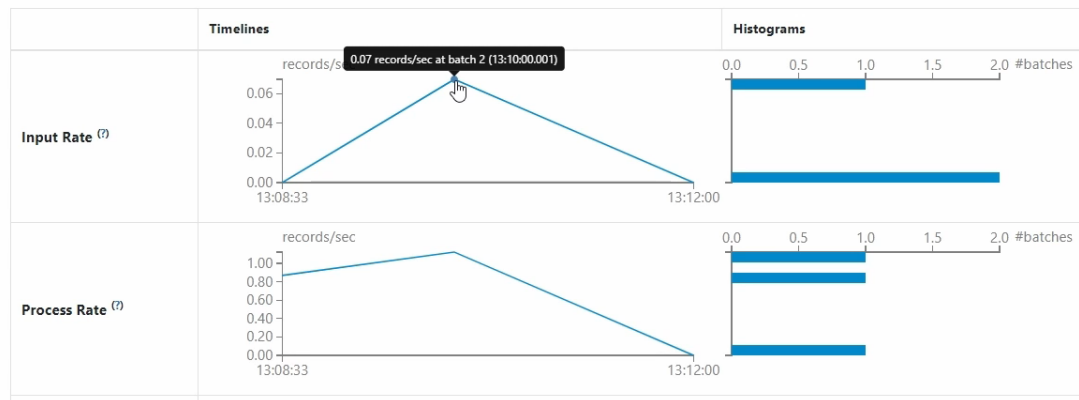
Streaming Query Statistics

Running batches for **4 minutes 58 seconds** since **2023/12/06 13:08:33** (4 completed batches)

Name: ProcessingTime

Id: a279ada4-dcc7-4770-9ceb-72f41d02293f

RunId: 6e895f53-48ff-4cfe-9c8c-68896bfcbb36



At 13:08 initial processing occurred, even though data is available only after 2 min next micro batch will be processed.

Cmd 12

03. Trigger - availablenow

```

1 WriteStream = ( df.writeStream
2   .option('checkpointLocation',f'{source_dir}/AppendCheckpoint')
3   .outputMode("append")
4   .trigger(availableNow=True)
5   .queryName('AvailableNow')
6   .toTable("stream.AppendTable"))

```

▶ (1) Spark Jobs

▶ AvailableNow (id: 9150eaac-ace9-4fd7-a9dc-07c5dc687ac4) Last updated: 45 seconds ago

Control console

Is equivalent to incremental batch processing, whenever there is a new data available query will process and stop the streaming

Triggers

Triggers	Usage	Description
Unspecified (default)		will trigger the microbatch for every 500 ms or half a second
processingTime (Fixed Interval)	.trigger(processingTime='2 minutes')	You can set processing time or time interval for each execution .
availableNow (OneTime)	.trigger(availableNow = True)	consumes all available records from previous execution as an incremental batch
Continuous (experimental)	.trigger(continuous = '1 second')	For ~1ms latency

Types of streaming triggers

Autoloader

Autoloader

- Autoloader is an **optimized data ingestion tool** that incrementally and efficiently processes new data files as they arrive in the cloud storage built into the Databricks Lakehouse.
- Auto Loader incrementally and efficiently processes new data files as they arrive in cloud storage without any additional setup.
- Auto Loader can load data files from Cloud Storages without being vendor specific (AWS S3 , Azure ADLS , Google Cloud Storage, DBFS).
- Auto Loader can ingest JSON, CSV, PARQUET, AVRO, ORC, TEXT, and BINARYFILE file formats
- This Auto loader is beneficial when you are ingesting data into your lakehouse particularly into bronze layer as a streaming query.

```
df_str = (spark.readStream
    .format("cloudFiles")  ## This will tell the spark to use AutoLoader.
    .option("cloudFiles.format", "csv") ## Tells Autoloader to expect csv files
    .option('header', 'true')
    .schema(schema)
    .load(f'{source_dir}')
)
```

To use the autoloader format("cloudFiles") needs to be specified

AutoLoader

Cmd 4

```
1 source_dir = 'dbfs:/FileStore/streaming/'
```

Command took 0.14 seconds -- by shanmukhsattiraju999@gmail.com at 06/12/2023, 21:41:15 on new

Cmd 5

```
1 df = spark.readStream\
2   .format('cloudFiles')\
3   .option("cloudFiles.format","csv")\
4   .option("cloudFiles.schemaLocation",f'{source_dir}/schemaInfer')\
5   .option("cloudFiles.inferColumnTypes","true")\
6   .option('header','true')\
7   .load(source_dir)
```

df: pyspark.sql.dataframe.DataFrame

Country: string
Citizens: integer
_rescued_data: string

Command took 0.47 seconds -- by shanmukhsattiraju999@gmail.com at 06/12/2023, 21:47:32 on new

To get the correct data types, schema must be inferred.

```
1 %sql
2
3 SELECT *
4 FROM JSON.`dbfs:/FileStore/streaming/schemaInfer/_schemas/0`
```

(2) Spark Jobs

_sqldf: pyspark.sql.dataframe.DataFrame = [_corrupt_record: string, dataSchemaJson: string ... 1 more field]

Table

	_corrupt_record	dataSchemaJson	partitionSchemaJson
1	v1	null	null
2	null	{ "type": "struct", "fields": [{ "name": "Country", "type": "string", "nullable": true, "metadata": {} }, { "name": "Citizens", "type": "integer", "nullable": true, "metadata": {} }] }	{ "type": "struct", "fields": [] }

Schema can be referred through the path, rescued data is the additional field created.

Schema Hints

For explicitly defining the data type of a column

SchemaHints

Cnd 10

```
1 df = spark.readStream\  
2     .format('cloudFiles')\  
3     .option("cloudFiles.format","csv")\  
4     .option("cloudFiles.schemaLocation",f'{source_dir}/schemaInfer')\  
5     .option("cloudFiles.inferColumnTypes","true")\  
6     .option('cloudFiles.schemaHints',"Citizens LONG")\  
7     .option('header','true')\  
8     .load(source_dir)
```

▼ df: pyspark.sql.dataframe.DataFrame

```
Country: string  
Citizens: long  
_rescued_data: string
```

Command took 0.32 seconds -- by sharmukhsettiraju999@gmail.com at 06/12/2023, 21:51:01 on new

Cnd 11

For explicitly defining the data type of a column

Schema Evolution

- Schema evolution is the process of managing changes in data schema as it evolves over time, often due to updates in software or changing business requirements, which can cause schema drift
- Ways to handle schema changes
 - Fail the stream
 - Manually change the existing schema
 - Evolve automatically with change in schema

There are 4 schema evolution types

1. AddNewColumns

Stream will fail, if new column is added and existing columns do not evolve.

2. FailOnNewColumns

Stream fails, stream does not restart until schema is updated or offending data is removed.

3. Rescue

Schema never evolved; stream doesn't change due to schema changes. All new columns are recorded in the rescued data.

4. None

Ignore new columns, (doesn't evolve the schema, new columns are ignored, data is not rescued unless the rescuedDataColumn option is set. Stream doesn't fail due to schema changes)

Using the **rescue type schema evolution**

Schema Evolution

Cmd 14

```
1 df = spark.readStream\  
2     .format('cloudFiles')\  
3     .option("cloudFiles.format","csv")\  
4     .option("cloudFiles.schemaLocation",f'{source_dir}/schemaInfer')\  
5     .option('cloudFiles.schemaEvolutionMode','rescue')\  
6     .option('rescuedDataColumn','_rescued_data')\  
7     .option("cloudFiles.inferColumnTypes","true")\  
8     .option('cloudFiles.schemaHints',"Citizens LONG")\  
9     .option('header','true')\  
10    .load(source_dir)
```

df: pyspark.sql.dataframe.DataFrame

Country: string
Citizens: long
_rescued_data: string

Table +

	Country	Citizens	_rescued_data
1	India	10	{"Month":"November","Year":"2023","_file_path":"dbfs/FileStore/streaming/Countries_newcolumn1.csv"}
2	USA	10	{"Month":"November","Year":"2023","_file_path":"dbfs/FileStore/streaming/Countries_newcolumn1.csv"}
3	China	20	{"Month":"November","Year":"2023","_file_path":"dbfs/FileStore/streaming/Countries_newcolumn1.csv"}
4	Brazil	10	{"Month":"November","Year":"2023","_file_path":"dbfs/FileStore/streaming/Countries_newcolumn1.csv"}
5	India	10	null
6	USA	5	null
7	China	10	null

Extra columns that are added handled by the rescued_data column

02 - addNewColumns - Default

Cmd 19

```
1 df = spark.readStream\
2   .format('cloudFiles')\
3   .option("cloudFiles.format","csv")\
4   .option("cloudFiles.schemaLocation",f'{source_dir}/schemaInfer')\
5   .option("cloudFiles.inferColumnTypes","true")\
6   .option("cloudFiles.schemaHints","Citizens LONG")\
7   .option('header','true')\
8   .load(source_dir)
```

df: pyspark.sql.dataframe.DataFrame = [Country: string, Citizens: long ... 1 more field]

Command took 0.26 seconds -- by shanmukhsettiraju999@gmail.com at 06/12/2023, 22:46:17 on new

Cmd 20

```
1 display(df)
```

(1) Spark Jobs

display_query_7 (id: 60120d56-4420-4ca8-b356-625185f2b644) Last updated: 5 seconds ago

org.apache.spark.sql.catalyst.util.UnknownFieldException: Encountered unknown field(s) during parsing: [Year, Month] in CSV file: dbfs://fileStore/streaming/Countries_newcolumn1.csv

Command complete

In the addNewColumns(default) type, after adding a file with extra columns, dataframe will be executed, however when we do a display it will fail.

```
1 display(df)
```

Cancel

(1) Spark Jobs

display_query_8 (id: fb45b5ee-c152-4d90-8227-75c77ce0a14a) Last updated: 10 seconds ago

	Country	Citizens	Year	Month	_rescued_data
1	India	10	2023	November	null
2	USA	10	2023	November	null
3	China	20	2023	November	null
4	Brazil	10	2023	November	null
5	India	10	null	null	null
6	USA	5	null	null	null
7	China	10	null	null	null

To fix this need to rerun the data frame and display to see the data with evolved schema.

03- failOnNewColumns

Cmd 24

```
1 dbutils.fs.ls('dbfs:/FileStore/streaming/schemaInfer/_schemas/')
```

```
Out[33]: [FileInfo(path='dbfs:/FileStore/streaming/schemaInfer/_schemas/0', name='0', size=274, modificationTime=1701879153000)]
```

Command took 0.11 seconds -- by shanmukhsattiraju999@gmail.com at 06/12/2023, 22:51:58 on new

Cmd 25

```
1 df = spark.readStream\  
2     .format('cloudFiles')\  
3     .option("cloudFiles.format","csv")\  
4     .option("cloudFiles.schemaLocation",f'{source_dir}/schemaInfer')\  
5     .option("cloudFiles.schemaEvolutionMode",'failOnNewColumns')\  
6     .option("cloudFiles.inferColumnTypes","true")\  
7     .option("cloudFiles.schemaHints","Citizens LONG")\  
8     .option('header','true')\  
9     .load(source_dir)
```

▶ df: pyspark.sql.dataframe.DataFrame = [Country: string, Citizens: long ... 1 more field]

Command took 0.34 seconds -- by shanmukhsattiraju999@gmail.com at 06/12/2023, 22:52:47 on new

Cmd 26

FailOnNewcolumns schema evolution mode, when a new file is added, display will fail, to fix either schema must be redefined or data should be changed.

None

Cmd 29

```
1 df = spark.readStream\  
2     .format('cloudFiles')\  
3     .option("cloudFiles.format","csv")\  
4     .option("cloudFiles.schemaLocation",f'{source_dir}/schemaInfer')\  
5     .option("cloudFiles.schemaEvolutionMode",'none')\  
6     .option("cloudFiles.inferColumnTypes","true")\  
7     .option("cloudFiles.schemaHints","Citizens LONG")\  
8     .option('header','true')\  
9     .load(source_dir)
```

▶ df: pyspark.sql.dataframe.DataFrame = [Country: string, Citizens: long]

Command took 0.27 seconds -- by shanmukhsattiraju999@gmail.com at 06/12/2023, 22:54:44 on new

Cmd 30

```
1 display(df)
```

Cancel ==>

▶ (1) Spark Jobs

▶ display_query_10 (id: 573f1cbf-356e-4e9e-bf94-23278455c8e7) Last updated: 10 seconds ago

None type of schema evolution ignores any changes in the schema, add the data to the existing schema.

